

# Churn Modelling Data

```
In [1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [109...]: data = pd.read_excel("P3- Churn-Modelling Data.xlsx")  
data.head(2)
```

```
Out[109...]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Ten
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	

```
In [3]: data.isnull().sum()
```

```
Out[3]:
```

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
churned	0

dtype: int64

```
In [4]: print(data.dtypes)
```

RowNumber	int64
CustomerId	int64
Surname	object
CreditScore	int64
Geography	object
Gender	object
Age	int64
Tenure	int64
Balance	float64
NumOfProducts	int64
HasCrCard	int64
IsActiveMember	int64
EstimatedSalary	float64
churned	int64

dtype: object

```
In [5]: df = pd.DataFrame(data)
```

# 1. Customer Demographics:

- What is the distribution of customers across different age groups? - Analyze the gender distribution of customers.


## What is the distribution of customers across different age groups?

```
In [6]: df['Age Group'] = pd.cut(df['Age'], bins=range(0, 101, 10), right=False)

plt.figure(figsize=(9, 6))
barplot = sns.countplot(x='Age Group', data=df, palette='gray',)

for bar in barplot.patches:      # Adding count annotations on top of each bar
    bar_height = bar.get_height()
    barplot.annotate(format(bar_height, '.0f'),
                      (bar.get_x() + bar.get_width() / 2, bar_height),
                      ha='center', va='center',
                      size=10, xytext=(0, 8),
                      textcoords='offset points')

plt.xlabel('Age Group ----->')
plt.ylabel('No. of customers ----->')
plt.title('Distribution of Customers in Each Age Group')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()
```



```
In [7]: freq_customer = df["Age"].value_counts()
freq_customer
```

```
Out[7]: 37    478
38    477
35    474
36    456
34    447
...
92     2
82     1
88     1
85     1
83     1
Name: Age, Length: 70, dtype: int64
```

```
In [8]: # Convert freq_customer Series to a DataFrame with appropriate column name
freq_customer_df = freq_customer.reset_index()
freq_customer_df.columns = ['Age', 'Count']

# Export to Excel
excel_file_path = 'frequency_counts.xlsx' # Replace with your desired file
freq_customer_df.to_excel(excel_file_path, index=False)

print(f"Frequency counts exported to {excel_file_path}")
```

Frequency counts exported to frequency\_counts.xlsx


## Analyze the gender distribution of customers.

```
In [9]: # df["Gender"].unique()
Gender_distribution = df["Gender"].value_counts()
Gender_distribution
```

```
Out[9]: Male      5457
Female    4543
Name: Gender, dtype: int64
```

```
In [10]: plt.figure(figsize=(8, 4))
plt.pie(Gender_distribution, labels=['Male', 'Female'], autopct='%1.2f%%')
plt.title('Gender distribution of customers\n')
plt.axis('equal')
plt.show()
```

Gender distribution of customers



## 2. Churn Analysis:

- What percentage of customers have churned?
- What are the main reasons for customer churn?
- Identify any patterns or trends among customers who have churned.

### - What percentage of customers have churned?

```
In [11]: data.head(2)
```

```
Out[11]:   RowNumber CustomerId Surname CreditScore Geography Gender Age Ten
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Ten
0	1	Hargrave	619	France	Female	42	
1	2	Hill	608	Spain	Female	41	


```
In [12]: churned = df['churned'].value_counts()
churned
```

```
Out[12]: 0    7963
1    2037
Name: churned, dtype: int64
```

```
In [13]: plt.figure(figsize=(8, 4))
plt.pie(churned, labels=['Not churned', 'Churned'], autopct='%1.2f%%', startangle=90)
plt.title('Percentage of customers churned')
```

```
plt.axis('equal')
plt.show()
```

Percentage of customers churned




## - What are the main reasons for customer churn?

```
In [14]: correlation = data.corr()
```

```
/var/folders/0c/mt12df7d7rj2s2x_hbbf4mvw0000gn/T/ipykernel_63391/307749493.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.
correlation = data.corr()
```

```
In [15]: plt.figure(figsize=(6, 4))
sns.heatmap(correlation, annot=True, fmt=".2f", cmap="coolwarm")
plt.title("Correlation Matrix")
plt.show()
```



```
In [16]: reasons_for_churn = "'Age' and 'Balance' are the main reasons for customer churn."
print(reasons_for_churn)
```


'Age' and 'Balance' are the main reasons for customer churn.

- Identify any patterns or trends among customers who have churned.


```
In [17]: # Analyze age distribution among churned customers
plt.figure(figsize=(6, 4))
sns.histplot(data=df[df['churned'] == 1], x='Age', bins=30, kde=True, col
plt.title('Age Distribution of Churned Customers')
plt.xlabel('Age ----->')
plt.ylabel('Total Customer----->')
plt.show()

# Analyze balance distribution among churned customers
plt.figure(figsize=(6,4))
sns.histplot(data=df[df['churned'] == 1], x='Balance', bins=30, kde=True,
plt.title('Balance Distribution of Churned Customers')
plt.xlabel('Balance ----->')
plt.ylabel('Total Customer----->')
plt.show()
```

### Age Distribution of Churned Customers



### Balance Distribution of Churned Customers



## 3. Product Usage:

- What are the most commonly used products or services? - Analyze the usage patterns of different customer segments.#

In [18]: `data.head(2)`

Out[18]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Ten
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	

In [19]: `print(df['Age Group'].unique())`

```
[[40, 50), [30, 40), [50, 60), [20, 30), [60, 70), [10, 20), [70, 80), [80, 90), [90, 100)]
Categories (10, interval[int64, left]): [[0, 10) < [10, 20) < [20, 30) < [30, 40) ... [60, 70) < [70, 80) < [80, 90) < [90, 100)]
```

In [20]: `print(df['Age Group'].value_counts())`

[30, 40)	4346
[40, 50)	2618
[20, 30)	1592
[50, 60)	869
[60, 70)	375
[70, 80)	136
[10, 20)	49
[80, 90)	13
[90, 100)	2
[0, 10)	0

Name: Age Group, dtype: int64

In [ ]:

- What are the most commonly used products or services?

In [21]: `print(df["NumOfProducts"].value_counts()),  
print(df["HasCrCard"].value_counts()),  
print(df["IsActiveMember"].value_counts())`

1	5084
2	4590
3	266
4	60

Name: NumOfProducts, dtype: int64

1	7055
0	2945

Name: HasCrCard, dtype: int64

1	5151
0	4849

Name: IsActiveMember, dtype: int64

- Analyze the usage patterns of different customer segments.

In [25]: `df.head(3)`

Out [25]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Ten
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	

In [65]:

```
# Segmentation
df['Age Group'] = pd.cut(df['Age'], bins=range(0, 101, 10), right=False)
df['CreditScoreGroup'] = pd.cut(data['CreditScore'], bins=range(0, 1001, 10))

# Analysis
age_group_analysis = data.groupby('Age Group').mean()
gender_analysis = data.groupby('Gender').mean()
geography_analysis = data.groupby('Geography').mean()

# Print insights
# print("Age Group Analysis:\n", age_group_analysis)
# print("Gender Analysis:\n", gender_analysis)
# print("Geography Analysis:\n", geography_analysis)
```

```
/var/folders/0c/mt12df7d7rj2s2x_hbbf4mvw0000gn/T/ipykernel_63391/2666242254.py:6: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is deprecated. In a future version, numeric_only will default to False. Either specify numeric_only or select only columns which should be valid for the function.
    age_group_analysis = data.groupby('Age Group').mean()
/var/folders/0c/mt12df7d7rj2s2x_hbbf4mvw0000gn/T/ipykernel_63391/2666242254.py:7: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is deprecated. In a future version, numeric_only will default to False. Either specify numeric_only or select only columns which should be valid for the function.
    gender_analysis = data.groupby('Gender').mean()
/var/folders/0c/mt12df7d7rj2s2x_hbbf4mvw0000gn/T/ipykernel_63391/2666242254.py:8: FutureWarning: The default value of numeric_only in DataFrameGroupBy.mean is deprecated. In a future version, numeric_only will default to False. Either specify numeric_only or select only columns which should be valid for the function.
    geography_analysis = data.groupby('Geography').mean()
```

In [32]:

```
credit_card_counts = df[df['HasCrCard'] == 1].groupby('Age Group').size()

print(credit_card_counts)
```

```
Age Group
[0, 10)      0
[10, 20)     36
[20, 30)    1136
[30, 40)    3092
[40, 50)    1830
[50, 60)     588
[60, 70)     263
[70, 80)     100
[80, 90)      9
[90, 100)     1
dtype: int64
```

```
In [76]: Active_member_counts = df[df['IsActiveMember'] == 1].groupby('Age Group')
print(Active_member_counts)
```


```
Age Group
[0, 10)      0
[10, 20)     28
[20, 30)    809
[30, 40)   2181
[40, 50)   1222
[50, 60)     496
[60, 70)     281
[70, 80)     120
[80, 90)      12
[90, 100)     2
dtype: int64
```

```
In [77]: # Active_member_counts = df[df['Geography'] == 1].groupby('Age Group').size()
# print(Active_member_counts)
Geography_region = df['Geography'].unique()
Geography_region
```


```
Out[77]: array(['France', 'Spain', 'Germany'], dtype=object)
```

```
In [78]: plt.figure(figsize = (6,4))
bar = credit_card_counts.plot(kind='bar', edgecolor='black')
for bar in barplot.patches:          # Adding count annotations on top of each bar
    bar_height = bar.get_height()
    barplot.annotate(format(bar_height, '.0f'), (bar.get_x() + bar.get_width()/2, bar.get_y() + bar_height))
plt.title('Number of Customers with Credit Card by Age Group')
plt.xlabel('Age Group ----->')
plt.ylabel('Number of Customers ----->')
plt.show()
```

### Number of Customers with Credit Card by Age Group



```
In [80]: plt.figure(figsize = (6,4))
bar = Active_member_counts.plot(kind='bar', edgecolor='black')
for bar in barplot.patches:          # Adding count annotations on top of each bar
    bar_height = bar.get_height()
    barplot.annotate(format(bar_height, '.0f'), (bar.get_x() + bar.get_width()/2, bar.get_y() + bar.get_height()))
plt.title('Number of Customers which are Active_member by Age Group')
plt.xlabel('Age Group ----->')
plt.ylabel('Number of Customers ----->')
plt.show()
```



## 4. Financial Analysis:

- What is the average account balance of customers? - Compare the financial characteristics of churned vs. non-churned customers.

```
In [ ]: data.head(2)
```

```
In [ ]: data.describe()
```

- What is the average account balance of customers?

```
In [ ]: df['Balance'].mean()
```

- Compare the financial characteristics of churned vs. non-churned customers.

```
In [ ]: total_churned_customer = df.groupby('churned')['Balance'].sum().reset_index()
total_churned_customer
```

```
In [ ]: total_churned_customer1 = df.groupby('churned')['EstimatedSalary'].sum()
total_churned_customer1
```

```
In [ ]: plt.figure(figsize=(8, 4))

# Subplot 1: Total Balance Distribution
plt.subplot(1, 2, 1)
plt.pie(total_churned_customer['Balance'], labels=['Not Churned', 'Churned'])
plt.title('Total Balance Distribution')
plt.axis('equal')

# Subplot 2: Total Estimated Salary Distribution
plt.subplot(1, 2, 2)
plt.pie(total_churned_customer1['EstimatedSalary'], labels=['Not Churned', 'Churned'])
plt.title('Total Estimated Salary Distribution')
plt.axis('equal')

# Adjust layout to prevent overlapping
plt.tight_layout()

plt.show()
```

## 5. Predictive Modeling:

- Which factors are the most significant predictors of customer churn? - Develop a predictive model to identify at-risk customers.

```
In [ ]: data.head(2)
```

```
In [ ]: # df
```

- Which factors are the most significant predictors of customer churn?**

```
In [ ]: # Age and Balance is highly correlated so these are most significant pred
```

- Develop a predictive model to identify at-risk customers**

```
In [86]: data.head(2)
```

```
Out[86]:   RowNumber CustomerId Surname CreditScore Geography Gender Age Ten
```

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Ten
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41

```
In [113]: # sns.heatmap(data , annot = True, cmap = 'coolwarm')

In [ ]:

In [ ]:

In [87]: #MODEL

In [88]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

In [89]: X = data[['Age', 'Balance']]
y = data['churned']

In [90]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In [106]: X_test
```

Out[106]:

	Age	Balance
6252	32	96709.07
4684	43	0.00
1731	44	0.00
4742	59	119152.10
4521	27	124995.98
...	...	...
6412	53	98268.84
8285	25	0.00
7853	47	0.00
1095	29	0.00
6929	39	115341.19

2000 rows × 2 columns

```
In [91]: model = LogisticRegression(max_iter = 1000)
model.fit(X_train, y_train)

Out[91]:
```

LogisticRegression

LogisticRegression(max\_iter=1000)


```
In [92]: y_pred = model.predict(X_test)
y_pred
```

```
Out[92]: array([0, 0, 0, ..., 0, 0, 0])
```

```
In [93]: print(f'Accuracy: {accuracy_score(y_test, y_pred)}')
print('Confusion Matrix:')
confusion_matrix = (confusion_matrix(y_test, y_pred))

plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix, annot=True, fmt="d", cmap="Greens")
plt.xlabel("Predicted labels")
plt.ylabel("True labels")
plt.title("Confusion Matrix")
plt.show()
```

Accuracy: 0.8035  
 Confusion Matrix:



```
In [94]: accuracy = accuracy_score(y_pred, y_test)
accuracy
```

Out[94]: 0.8035

```
In [95]: print('Classification Report:')
print(classification_report(y_test, y_pred))
```

Classification Report:				
	precision	recall	f1-score	support
0	0.80	1.00	0.89	1607
1	0.00	0.00	0.00	393
accuracy			0.80	2000
macro avg	0.40	0.50	0.45	2000
weighted avg	0.65	0.80	0.72	2000

```
/Users/rakeshmanawat/anaconda3/lib/python3.11/site-packages/sklearn/metric
s/_classification.py:1469: UndefinedMetricWarning: Precision and F-score a
re ill-defined and being set to 0.0 in labels with no predicted samples. U
se `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/Users/rakeshmanawat/anaconda3/lib/python3.11/site-packages/sklearn/metric
s/_classification.py:1469: UndefinedMetricWarning: Precision and F-score a
re ill-defined and being set to 0.0 in labels with no predicted samples. U
se `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/Users/rakeshmanawat/anaconda3/lib/python3.11/site-packages/sklearn/metric
s/_classification.py:1469: UndefinedMetricWarning: Precision and F-score a
re ill-defined and being set to 0.0 in labels with no predicted samples. U
se `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

In [96]: `#random forest`

```
In [97]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.metrics import r2_score

# Initialize the model
model = RandomForestClassifier(random_state=42)

# Train the model
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)


print(f'Accuracy: {accuracy}')
```

Accuracy: 0.768

In [98]: `print('Confusion Matrix:')`  
`cm = confusion_matrix(y_test, y_pred)`

```
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Greens")
plt.xlabel("Predicted labels")
plt.ylabel("True labels")
plt.title("Confusion Matrix")
plt.show()
```

Confusion Matrix:



```
In [99]: report = classification_report(y_test, y_pred)
print(report)
```

	precision	recall	f1-score	support
0	0.84	0.88	0.86	1607
1	0.39	0.33	0.36	393
accuracy			0.77	2000
macro avg	0.62	0.60	0.61	2000
weighted avg	0.75	0.77	0.76	2000

```
In [ ]:
```

linear regression

```
In [100]: from sklearn.linear_model import LogisticRegression
```

```
In [101]: model1 = LogisticRegression()
model1.fit(X, y)
```

```
Out[101]: ▾ LogisticRegression
          LogisticRegression()
```

```
In [102]: y_pred = model1.predict(X)
y_pred
```

```
Out[102]: array([0, 0, 1, ..., 0, 1, 0])
```

```
In [103]: r2 = r2_score(y, y_pred)
print(f'R-squared: {r2}') #coefficient of determination
```

R-squared: 0.4710440056246886

```
In [ ]: 
```

```
In [ ]: 
```

```
In [ ]: 
```

```
In [104... # SAVE THE MODEL
```

```
In [107... import numpy as np
import pickle
```

```
In [108... from sklearn.linear_model import LogisticRegression

# Example: Training a model
model1 = LogisticRegression()
model1.fit(X_train, y_train)

# Saving the model to a file
with open('model1.pkl', 'wb') as file:
    pickle.dump(model, file)

# Loading the model from the file
with open('model1.pkl', 'rb') as file:
    loaded_model = pickle.load(file)

# Using the loaded model to make predictions
predictions = loaded_model.predict(X_test)
predictions
```

```
Out[108... array([0, 0, 0, ..., 0, 0, 1])
```

```
In [ ]: # predictions = loaded_model.predict([[40,5000]])    #correct shape is 2D
# print(predictions)
```

```
In [114... # Predict on the test set
y_pred = model1.predict(X_test)

# Evaluate the model

accuracy = accuracy_score(y_test, y_pred)


print(f'Accuracy: {accuracy}')
```

```
Accuracy: 0.8035
```

```
In [117... print('Confusion Matrix:')
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(4, 2))
sns.heatmap(cm, annot=True, fmt="d", cmap="Greens")
plt.xlabel("Predicted labels")
plt.ylabel("True labels")
plt.title("Confusion Matrix")
plt.show()
```

Confusion Matrix:



```
In [116]: report1 = classification_report(y_test, y_pred)
print(report1)
```

	precision	recall	f1-score	support
0	0.80	1.00	0.89	1607
1	0.00	0.00	0.00	393
accuracy			0.80	2000
macro avg	0.40	0.50	0.45	2000
weighted avg	0.65	0.80	0.72	2000

```
/Users/rakeshmanawat/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/Users/rakeshmanawat/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/Users/rakeshmanawat/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

```
In [ ]:
```